


Logistic Regression Classification

Lecture 10 – HCCDA-AI

Classification

- Predicts discrete-valued output.
- Classification is the task of predicting a discrete label (or class) for an input.
- Output is categorical: each input is assigned to one of the predefined classes.



Common Examples:


- **Email:** Spam / Not Spam?
- **Online Transactions:** Fraudulent (Yes / No)?
- **Tumor:** Malignant / Benign ?

$$y \in \{0, 1\}$$

0: “Negative Class” (e.g., benign tumor)

1: “Positive Class” (e.g., malignant tumor)

Classifying Tumor as Malignant or Benign:



Steps:


- Apply Linear Regression
- Pick a threshold
- For this dataset, linear regression seems to work reasonably well.

Threshold classifier output $h_{\theta}(x)$ at 0.5:

If $h_{\theta}(x) \geq 0.5$, predict “y = 1”

If $h_{\theta}(x) < 0.5$, predict “y = 0”

Classifying Tumor as Malignant or Benign:



- A single new example can break it
- Applying linear regression to a classification problem usually is not a great idea!!

Conclusion:

- We need a model designed to predict probabilities and handle classification boundaries more reliably.

Logistic Regression



Logistic Regression

- Logistic regression predicts the probability that the input belongs to **class 1**, the class of interest.
- It uses the **sigmoid** (logistic function) to convert linear output into probabilities in the range $[0, 1]$.

Sigmoid Function:

$$g(z) = \frac{1}{1 + e^{-z}}$$

- Converts linear output into probabilities.
- Helps determine class labels based on a threshold (e.g., $\geq 0.5 \rightarrow$ class 1, $< 0.5 \rightarrow$ class 0).



Logistic Regression

Hypothesis Representation


Logistic Regression Model

Want $0 \leq h_{\theta}(x) \leq 1$

$$h_{\theta}(x) = \theta^T x$$

Sigmoid function
Logistic function

$$h_{\theta}(x) = g(\theta^T x)$$
$$g(z) = \frac{1}{1 + e^{-z}}$$



Interpretation of Hypothesis Output

$h_{\theta}(x)$ = estimated probability that $y = 1$ on input x

Example: If $x = \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} = \begin{bmatrix} 1 \\ \text{tumorSize} \end{bmatrix}$

$$h_{\theta}(x) = 0.7$$

Tell patient that 70% chance of tumor being malignant

“probability that $y = 1$, given x ,
parameterized by θ ”

$$\begin{aligned} P(y = 0|x; \theta) + P(y = 1|x; \theta) &= 1 \\ P(y = 0|x; \theta) &= 1 - P(y = 1|x; \theta) \end{aligned}$$

Logistic Regression

Decision Boundary

Decision Boundary




- A decision boundary is a line, surface, or hypersurface that separates different classes in a feature space.
- It guides how data points are classified in a supervised learning algorithm.

Decision Boundary in Logistic Regression:


- Logistic regression aims to find a proper fit for the decision boundary to classify new data accurately.
- It is linear boundary in two dimensional space.

Key Insights:

- **Linear Decision Boundary:**
 - Suitable for problems where classes can be separated by a straight line.
- **Non Linear Boundaries:**
 - More complex models like decision trees, SVMs, or neural networks handle nonlinear separations.




Decision Boundary



$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$


Predict “ $y = 1$ “ if $-3 + x_1 + x_2 \geq 0$

Non-linear decision boundaries



$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2)$$

Predict “ $y = 1$ ” if $-1 + x_1^2 + x_2^2 \geq 0$



$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_1^2 x_2 + \theta_5 x_1^2 x_2^2 + \theta_6 x_1^3 x_2 + \dots)$$

Logistic Regression

Cost Function

Training set: $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$

m examples $x \in \begin{bmatrix} x_0 \\ x_1 \\ \dots \\ x_n \end{bmatrix} \quad x_0 = 1, y \in \{0, 1\}$

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

How to choose parameters θ ?

Cost Function

Linear Regression: $J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2$



$$\text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) = \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h_{\theta}(x^{(i)}), y^{(i)})^2$$

- Squared error leads to a **non-convex** cost function in logistic regression.
- Gradient descent may get stuck in local minima.

Cost Function


- The squared error cost function work will for linear regression, but if we use this particular cost function for logistic regression, this will be a non-convex function of parameter θ .
- For logistic regression:
$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$
- If put this sigmoid function in cost function. The $J(\theta)$ will look like figure a) with many local optimum.
- Figure a) is a non-convex function, if gradient descent is used, it is not guaranteed to converge to global optimum.



Logistic Regression Cost Function

Binary Cross Entropy Loss

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$



Cost = 0 if $y = 1, h_{\theta}(x) = 1$

But as $h_{\theta}(x) \rightarrow 0$

$\text{Cost} \rightarrow \infty$


Captures intuition that if $h_{\theta}(x) = 0$,
(predict $P(y = 1|x; \theta) = 0$), but $y = 1$,
we'll penalize learning algorithm by a very
large cost.

➤ Loss is lowest when $h_{\theta}(x^{(i)})$ predicts close to true label $y^{(i)}$

Logistic Regression Cost Function

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

If $y = 0$



Logistic Regression

Simplifies Cost Function and Gradient Descent

Logistic regression cost function

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)})$$

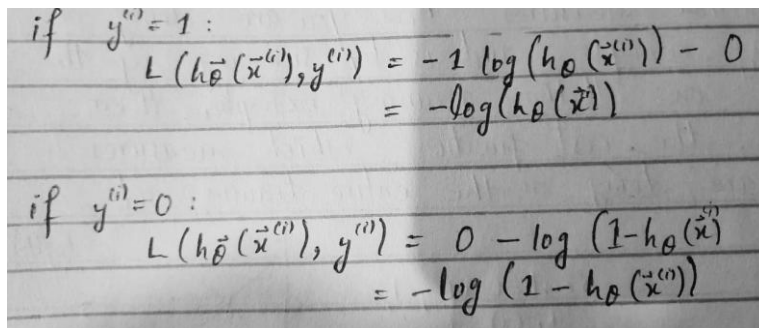
❖ Our overall loss function.

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

❖ For single training example

This equation is completely equivalent to the above more complex formula.

Note: $y = 0$ or 1 always



Handwritten derivation of the cost function for $y=1$ and $y=0$:

if $y^{(i)} = 1$:

$$L(h_{\theta}(\vec{x}^{(i)}), y^{(i)}) = -1 \log(h_{\theta}(\vec{x}^{(i)})) - 0 = -\log(h_{\theta}(\vec{x}^{(i)}))$$

if $y^{(i)} = 0$:

$$L(h_{\theta}(\vec{x}^{(i)}), y^{(i)}) = 0 - \log(1 - h_{\theta}(\vec{x}^{(i)})) = -\log(1 - h_{\theta}(\vec{x}^{(i)}))$$

Logistic regression cost function

$$\begin{aligned} J(\theta) &= \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) \\ &= -\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \right] \end{aligned}$$

To fit parameters : θ

$$\min_{\theta} J(\theta)$$

To make a prediction given new x :

$$\text{Output } h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

Gradient Descent

$$J(\theta) = -\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \right]$$

Want $\min_{\theta} J(\theta)$:

Repeat {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

}

(simultaneously update all θ_j)

Gradient Descent

$$J(\theta) = -\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \right]$$

Want $\min_{\theta} J(\theta)$:

Repeat {

$$\theta_j := \theta_j - \alpha \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

} (simultaneously update all θ_j)

Algorithm looks identical to linear regression!

Gradient of Binary Cross Entropy Loss Function

Derivative of Cost Function in Logistic Regression:
While implementing logistic regression gradient descent we need to compute:
Derivative of cost function with respect to θ_0 and θ_j .

$$\text{cost} = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1-y^{(i)}) \log(1-h_{\theta}(x^{(i)}))]$$

$$\rightarrow h_{\theta}(\vec{x}) = \frac{1}{1+e^{-(\theta_0 + \vec{\theta} \cdot \vec{x})}} = \sigma(\theta_0 + \vec{\theta} \cdot \vec{x})$$

$$= \sigma(z)$$

$$\rightarrow z = \theta_0 + \vec{\theta} \cdot \vec{x}$$

$\frac{\partial \text{Cost}}{\partial \theta_j}$? $\frac{\partial \text{Cost}}{\partial \theta_0}$?

Loss Function: $L = -[y^{(i)} \log(h_{\theta}(x^{(i)})) + (1-y^{(i)}) \log(1-h_{\theta}(x^{(i)}))]$

$$\vec{\theta} = \begin{bmatrix} \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix} \quad \frac{\partial L}{\partial \theta_1}, \frac{\partial L}{\partial \theta_2}, \dots \rightarrow \frac{\partial L}{\partial \vec{\theta}}$$

$\nabla \frac{\partial L}{\partial \vec{\theta}}$ By applying chain rule → For simplicity I write θ_j .

$$\frac{\partial L}{\partial \theta} = \frac{\partial L}{\partial h_{\theta}(x^{(i)})} \times \frac{\partial h_{\theta}(x^{(i)})}{\partial \theta}$$

$$= \frac{\partial L}{\partial h_{\theta}(x^{(i)})} \times \frac{\partial h_{\theta}(x^{(i)})}{\partial z} \times \frac{\partial z}{\partial \theta}$$

Let's find all the three derivatives one by one and then multiply it to get $\frac{\partial L}{\partial \theta}$

$$\Rightarrow \frac{\partial L}{\partial h_{\theta}(x^{(i)})} = \frac{\partial}{\partial h_{\theta}(x^{(i)})} - [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1-y^{(i)}) \log(1-h_{\theta}(x^{(i)}))]$$

$$= - \left[y^{(i)} \times \frac{1}{h_{\theta}(x^{(i)})} + (1-y^{(i)}) \times \frac{1}{1-h_{\theta}(x^{(i)})} \right] \because \frac{\partial \log x}{\partial x} = \frac{1}{x}$$

$$= \frac{-y^{(i)}}{h_{\theta}(x^{(i)})} - \frac{(1-y^{(i)})}{1-h_{\theta}(x^{(i)})}$$

→ y is our prediction, it will be constant

→ Also find derivative of $1-h_{\theta}(x^{(i)})$ which is $0-1=-1$

Gradient of Binary Cross Entropy Loss Function

so - will get cancelled.

$$\frac{\partial L}{\partial h_\theta(x^{(i)})} = \frac{-y}{h_\theta(x^{(i)})} + \frac{1-y}{1-h_\theta(x^{(i)})}$$

$$\Rightarrow \frac{\partial h_\theta(x^{(i)})}{\partial z} = \frac{\partial}{\partial z} \left(\frac{1}{1+e^{-z}} \right)$$

\therefore chain rule

$$= \frac{-1}{(1+e^{-z})^2} \cdot (1+e^{-z})$$

$$= \frac{-1}{(1+e^{-z})^2} \cdot (-e^{-z})$$

$$= \frac{e^{-z}}{(1+e^{-z})^2}$$

$\Rightarrow h_\theta(x^{(i)})^2 = \frac{1}{(1+e^{-z})^2}$

$\Rightarrow e^{-z} = \frac{1-h_\theta(x^{(i)})}{h_\theta(x^{(i)})}$

\Downarrow

so

$$= \frac{1-h_\theta(x^{(i)})}{h_\theta(x^{(i)})} \times (h_\theta(x^{(i)}))^2$$

$$h_\theta(x^{(i)}) = \frac{1}{1+e^{-z}}$$

$$= \frac{1-h_\theta(x^{(i)})}{h_\theta(x^{(i)})} \times h_\theta(x^{(i)}) \cdot h_\theta(x^{(i)})$$

$1+e^{-z} = \frac{1}{h_\theta(x^{(i)})}$

$e^{-z} = \frac{1-h_\theta(x^{(i)})}{h_\theta(x^{(i)})}$

$$= h_\theta(x^{(i)}) (1-h_\theta(x^{(i)}))$$

$$\Rightarrow \frac{\partial Z}{\partial \theta} = \frac{\partial}{\partial \theta} (\theta_0 + \theta_1 x)$$

$$= 0 + x$$

$$= x$$

Now

$$\frac{\partial L}{\partial \theta} = \left[\frac{-y}{h_\theta(x^{(i)})} + \frac{(1-y)}{(1-h_\theta(x^{(i)}))} \right] \times (h_\theta(x^{(i)})) (1-h_\theta(x^{(i)})) \times x$$

$$= \left[\frac{-y(1-h_\theta(x^{(i)})) + (1-y)(h_\theta(x^{(i)}))}{h_\theta(x^{(i)})(1-h_\theta(x^{(i)}))} \right] \cdot (h_\theta(x^{(i)})) (1-h_\theta(x^{(i)})) \times x$$

$$= \left[\frac{-y + y h_\theta(x^{(i)}) + h_\theta(x^{(i)}) - y h_\theta(x^{(i)})}{h_\theta(x^{(i)})(1-h_\theta(x^{(i)}))} \right] \cdot (h_\theta(x^{(i)})) (1-h_\theta(x^{(i)})) \times x$$

Gradient of Binary Cross Entropy Loss Function

$$\begin{aligned}
 \frac{\partial L}{\partial \theta} &= (-y + h_{\theta}(x^{(i)})) \cdot x \\
 \frac{\partial L}{\partial \theta} &= (h_{\theta}(x^{(i)}) - y) \cdot x \\
 \frac{\partial L}{\partial \vec{\theta}} &= (h_{\vec{\theta}}(x^{(i)}) - y) \cdot \vec{x}
 \end{aligned}
 \quad
 \left\{
 \begin{aligned}
 \frac{\partial L}{\partial \theta_1} &= (h_{\theta}(x^{(i)}) - y) \cdot x_1 \\
 \frac{\partial L}{\partial \theta_2} &= (h_{\theta}(x^{(i)}) - y) \cdot x_2 \\
 &\vdots \\
 \frac{\partial L}{\partial \theta_n} &= (h_{\theta}(x^{(i)}) - y) \cdot x_n
 \end{aligned}
 \right.$$

For θ_0 :

$$\begin{aligned}
 \frac{\partial L}{\partial \theta_0} &= \frac{\partial L}{\partial h_{\theta}(x^{(i)})} \times \frac{\partial h_{\theta}(x^{(i)})}{\partial z} \times \frac{\partial z}{\partial \theta_0} \\
 &= \left[\frac{-y}{h_{\theta}(x^{(i)})} + \frac{(1-y)}{(1-h_{\theta}(x^{(i)}))} \right] \times (h_{\theta}(x^{(i)}))(1-h_{\theta}(x^{(i)})) \cdot 1 \\
 &= (h_{\theta}(x^{(i)}) - y) \cdot 1 \\
 \frac{\partial L}{\partial \theta_0} &= h_{\theta}(x^{(i)}) - y
 \end{aligned}
 \quad
 \left\{
 \begin{aligned}
 z &= \theta_0 + \theta_1 x \\
 \frac{\partial z}{\partial \theta_0} &= 1
 \end{aligned}
 \right.$$

Logistic Regression

Multi-class classification: One-vs-all

Multiclass classification

- Multi-class classification is the process of assigning each data instance to **one** class among **three or more** possible categories. Unlike binary classification, where there are only two classes, multi-class classification deals with **multiple** distinct labels.

Examples


- Classifying emails as Spam, Promotions, or Primary.
- Identifying handwritten digits (0-9)
- Categorizing images into Dogs, Cats, and Birds.

Common Algorithms:

- Logistic Regression (One-vs-Rest, Softmax Regression)
- Decision Trees & Random Forest
- Support Vector Machines (SVM, One-vs-One strategy)
- Neural Networks (Deep Learning models like CNNs for image classification)

Loss Function

- Cross Entropy Loss (Softmax Loss) is commonly used to measure prediction accuracy.




Multiclass classification

Email foldering/tagging: Work, Friends, Family, Hobby


Medical diagrams: Not ill, Cold, Flu

Weather: Sunny, Cloudy, Rain, Snow

Binary classification:





Multi-class classification:



One-vs-all (one-vs-rest):

One vs all is a strategy used in multiclass classification where a separate binary classification model is trained for each class.






Class 1:  x_1

Class 2: 

Class 3: 

$$h_{\theta}^{(i)}(x) = P(y = i|x; \theta) \quad (i = 1, 2, 3)$$



Newton's Method

Newton's Method

Newton's Method is a second-order optimization algorithm that uses both the gradient and the Hessian matrix to find the optimal solution.

Key Points:

- Uses **second-order derivatives** (Hessian matrix) for optimization.
- Faster convergence than **Gradient Descent** but computationally expensive.
- Used in logistic regression, SVMs, neural networks, and nonlinear optimization.

Formula:

$$w = w - H^{-1} \nabla J(w)$$

where:

- H = Hessian matrix (second derivatives of the cost function)
- $J(w)$ = gradient of the cost function

Newton's Method vs. Gradient Descent

Feature	Newton's Method	Gradient Descent
Uses	Second-order derivatives	First-order derivatives
Convergence	Faster	Slower
Computational Cost	Higher (Hessian inversion)	Lower
Step Size Adjustment	Not required	Requires tuning
Suitable for Large Datasets	No	Yes

Applications:

- **Logistic Regression** (Newton-Raphson method)
- **Support Vector Machines (SVMs)**
- **Neural Network Training (L-BFGS method)**

Key Takeaway:

- Newton's Method is **faster but computationally expensive** compared to Gradient Descent. Suitable for small to medium-sized problems where quick convergence is needed.

Thank You