

Transfer Learning

Utilizing Pre-trained Model for Better Results

Lecture 17 – HCCDA-AI

Dr. Muhammad Sajjad

R.A: Kaleem Ullah

R.A: Imran Nawar

3 August 2025

1


Overview

- Challenges in training custom Deep Learning models
- Transfer Learning
- How Transfer Learning Works
- Types of Transfer Learning
- Pre-trained Models
- ImageNet
- ImageNet Competition
- AlexNet Architecture
- LeNet5
- VGG16/19 Architecture
- Problems with Very Deep Networks
- Residual Learning (RESNET Architecture)
- INCEPTION Architecture
- 1x1 Convolution
- EfficientNet
- MobileNet
- Applications of Transfer Learning

Challenges in Training Custom Deep Learning Models


Data:

- Deep learning models typically require a large amount of labeled data to learn effectively.
- Gathering and labeling this data can be time-consuming and expensive.
- Without sufficient data, the model may not generalize well to new, unseen examples, leading to poor performance.



Training Time:

- Training deep learning models can be computationally intensive and time-consuming.
- Depending on the complexity of the model architecture, size of the dataset, and available computational resources, training can take days, weeks, or even longer.
- Longer training times also increase the cost associated with experimentation and model development.



Challenges in Training Custom Deep Learning Models

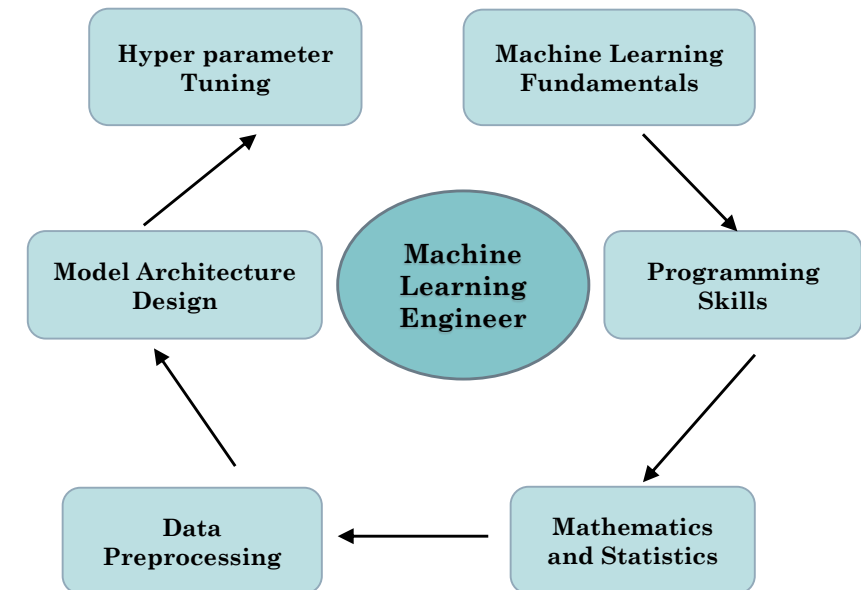
Computational Resources:

- Training deep learning models often requires significant computational resources, including powerful GPUs or even specialized hardware like TPUs.
- Not everyone has access to these resources, limiting the ability to train complex models effectively.



Expertise Requirement:

- Building and training custom deep learning models requires expertise in machine learning, deep learning, and software engineering.
- This expertise may not be readily available to everyone, especially those new to the field.



Solution

Transfer Learning

- A powerful technique in deep learning that allows us to reuse knowledge gained from solving one problem to tackle a different but related problem.
- Instead of training a model from scratch, we start with a pre-trained model and fine-tune it for the new task.
- It will not only speed up training considerably, but also requires significantly less training data.

Need of Transfer Learning?

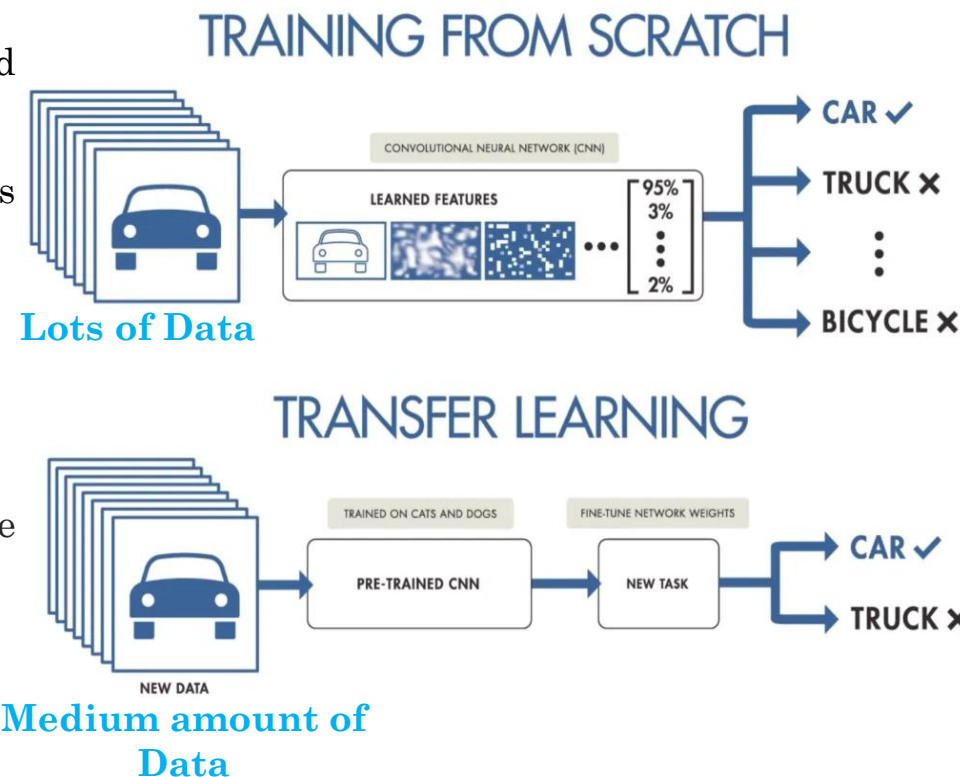
- If our dataset is really small.
- When training from scratch is computationally expensive.
- If our dataset is similar to pre-trained data then we have to only fine tuning our model it would save lot of time.

Advantages:

- Works with limited data
- Reduces training time
- Often improves model performance when pre-trained features are relevant.

Limitations

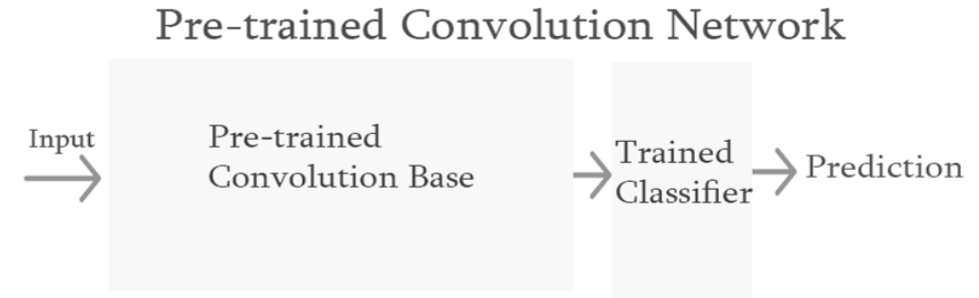
- Less effective if the dataset is very different from the pre-trained data.
- Fine-tuning may still be computationally intensive.



How Transfer Learning Works


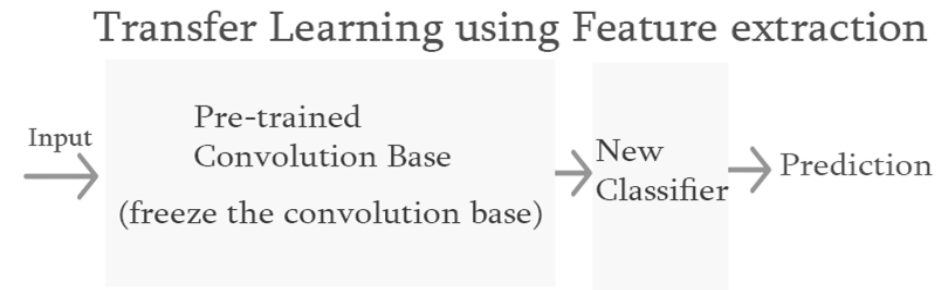
1. Pre-trained Model:

- Begin with a model that has already been trained on a large dataset for a specific task (e.g., image classification using ImageNet).
- This pre-trained model has learned useful features and patterns from the data.



2. Application to a New Task:

- Add new task-specific layers (e.g., an output layer) on top of the base model.
- Fine-tune the entire model on the new task using a smaller dataset.
- The base model's features are leveraged for the new task.



Types of Transfer Learning

1. Feature Extraction

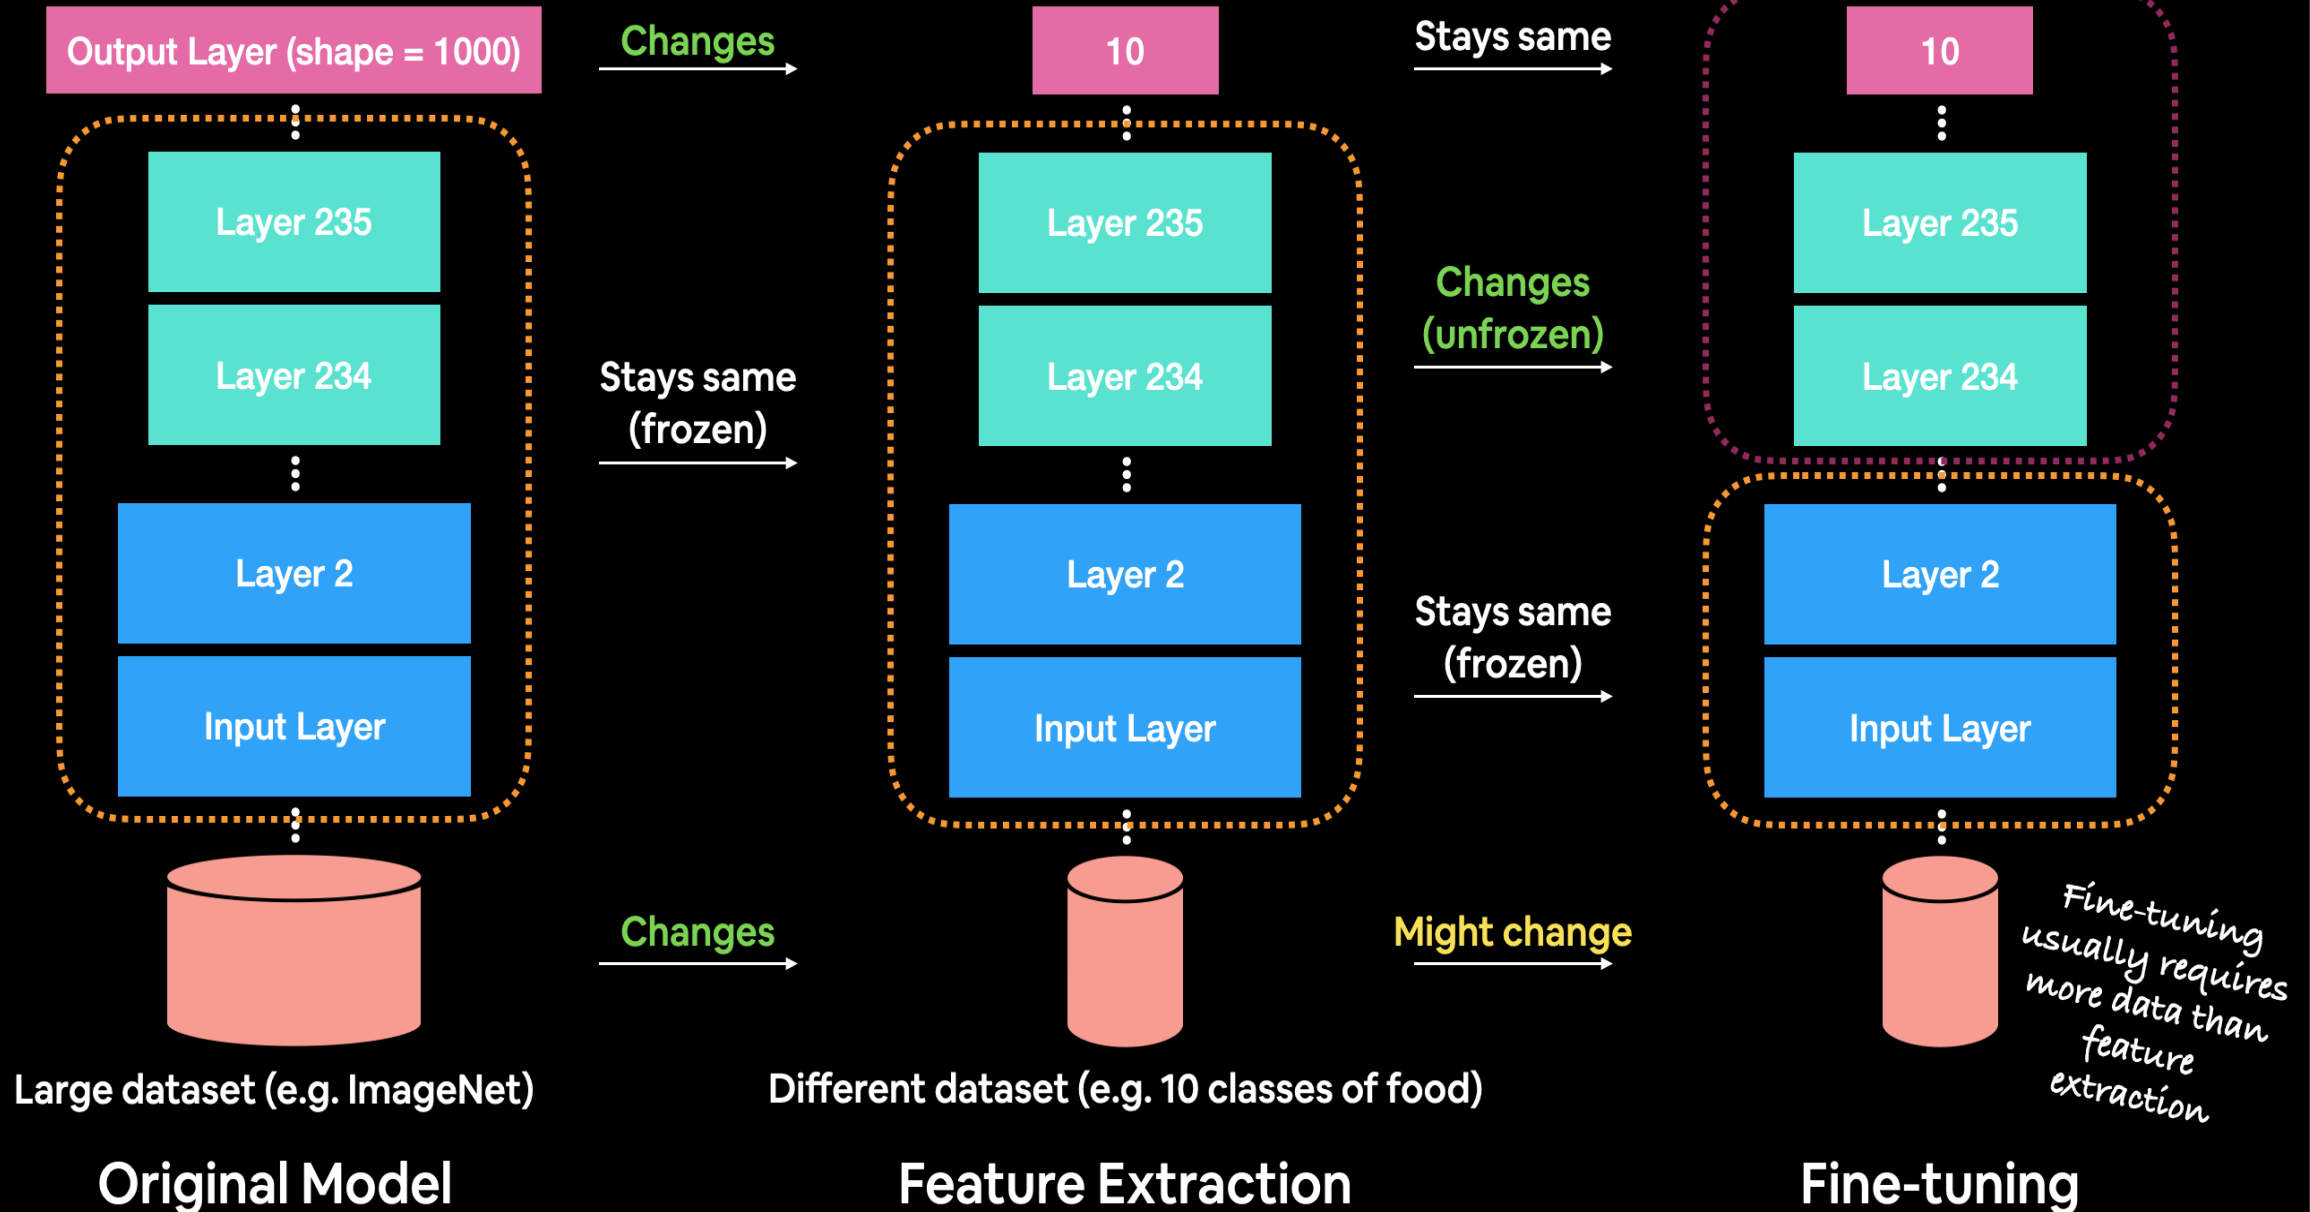
- Use a pre-trained model as a fixed feature extractor
- Extract features from the final convolutional layer (e.g., in a CNN)
- Feed these features into a new classifier (e.g., a fully connected layer or SVM)
- Commonly used when data for the new task is limited

2. Fine-Tuning:


- Unfreeze and fine-tune some or all layers of the pre-trained model
- Adjust the model's weights using the new dataset
- Useful when the new task is closely related to the original task

Types of Transfer Learning

Top layers get trained on new data




Transfer Learning



- ← - - Gradient flow
- ← No gradient flow
- Neurons from a pre-trained model
- Appended neurons

Fine Tuning



- ← - - Gradient flow
- Full pre-trained network
- Appended network

Pre-trained Models

Pre-trained Models

A pre-trained model is a deep learning model trained on a large benchmark dataset (e.g., ImageNet, COCO, BERT pre-training corpus).

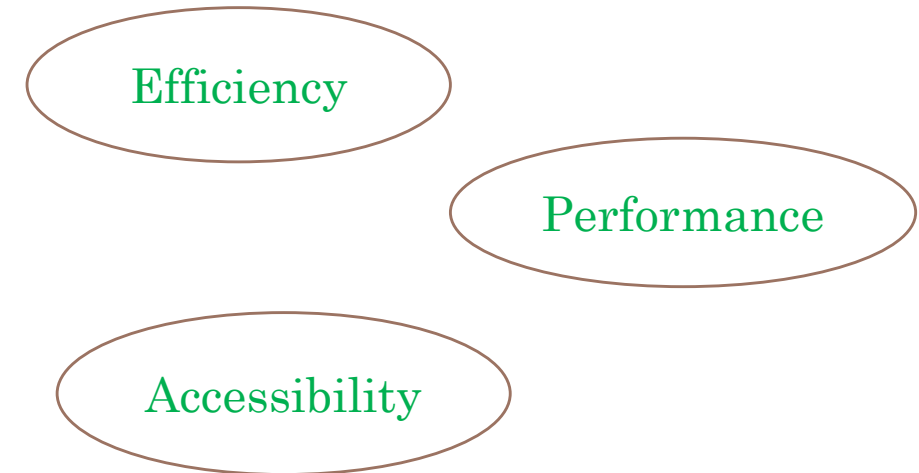
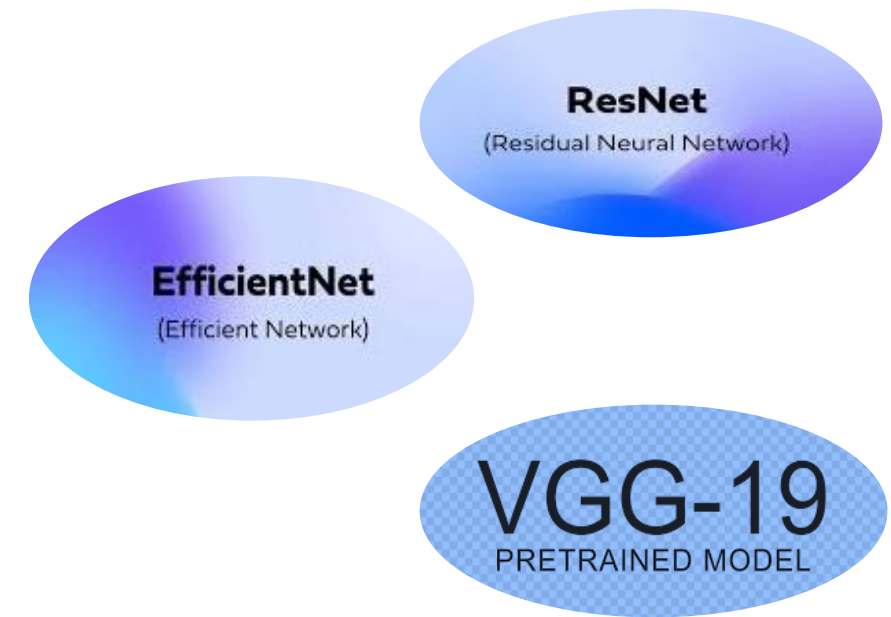
- These models serve as a starting point for new tasks by providing learned weights and feature representations.
- Typically, lower layers (which learn general features) are reused, while upper layers (task-specific) are fine-tuned or replaced.

Benefits of Using Pre-trained Models

- Reduce training time and computational cost
- Leverage knowledge from large datasets
- Improve accuracy, especially on small or similar datasets
- Enable faster prototyping and experimentation

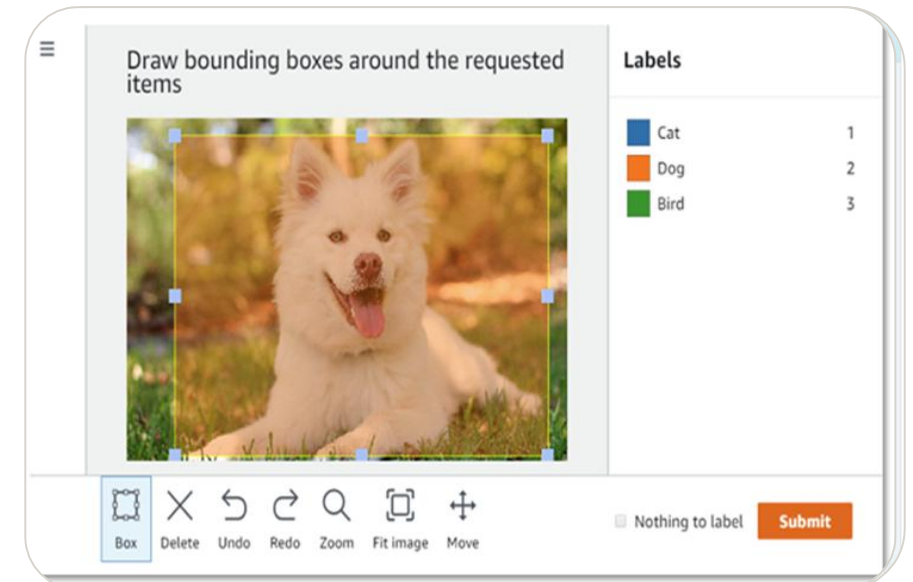

Example:

- **ResNet** trained on **ImageNet** (14 million images) can be fine-tuned to classify medical images with much less data.



ImageNet

- The ImageNet dataset contained over **14 million** labeled images.
- The dataset contains more than 20,000 categories.
- A prominent computer scientist, **Fei-Fei Li**, co-founded the ImageNet project in 2009 (initiated in 2006) along with other researchers. Their work on ImageNet significantly advanced computer vision research and deep learning development.
- There are **1 million** images with bounding box labelling as well for the purpose of **Object Localization** task, where the goal is to identify not only the object but also its precise location within the image.
- They used Amazon Mechanical Turk to help with the classification of images.
- **Amazon Mechanical Turk (MTurk)** is great for crowdsourcing tasks using images.



ImageNet Competition

- The ImageNet Large Scale Visual Recognition Challenge (ILSVRC) started in 2010 , and the goal was to Highlight the best model for classification to the research Community.
- The dataset which were used in the challenge was a **subset** of ImageNet which consists of around **1.2 million** images from **1000 classes**.
- Initially, participants used traditional **ML Algorithms**.
- The Error rate was **28%** in the first time.
- In **2011** the error rate reduces to **25%** using improved ML techniques.
- A revolution came in 2012 when **Geoffrey Hinton's** team participated with a CNN-based model, **AlexNet**.
- Its effective use of deep learning, ReLU activation, and GPU acceleration significantly boosted performance.
- AlexNet achieved an error rate of **15.3%**, a major improvement over previous methods.



Eight Years of Competitions

IMAGENET

2010-2017


10×
reduction of image
classification error

3×
improvement of
detection precision

Winning Models of the ImageNet Competition

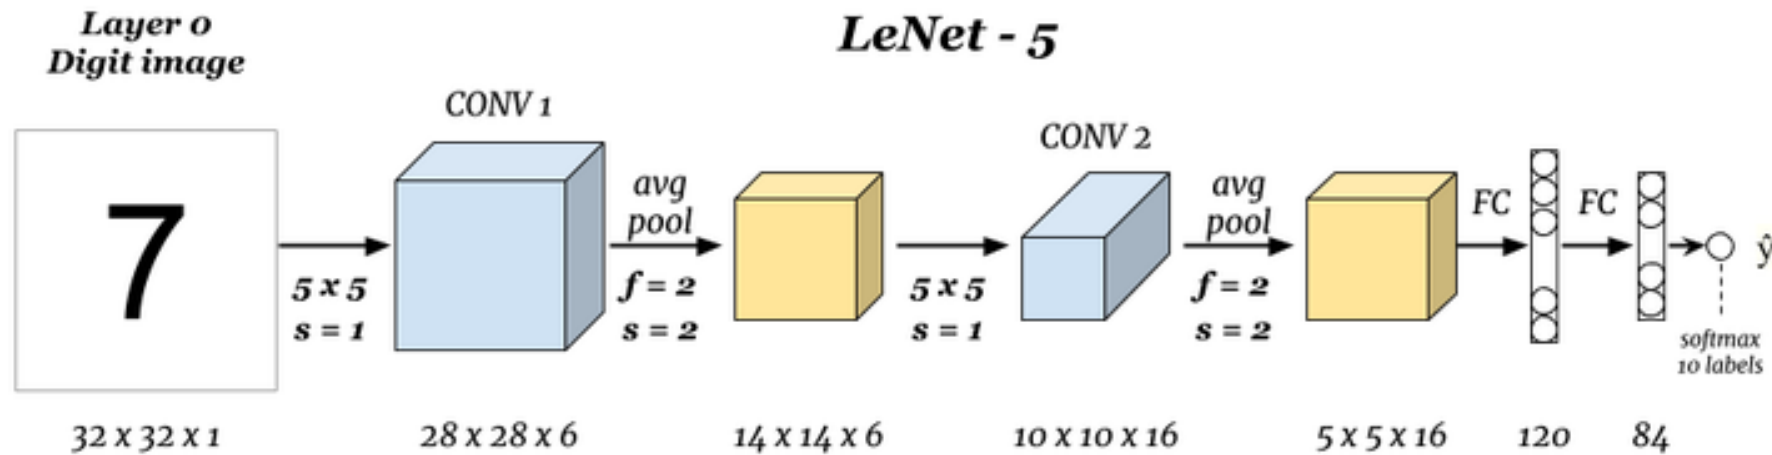
YEAR	WINNER	TOP 5 ERROR RATE %
2012	ALEXNET	15.3
2013	ZFNET	11.2
2014	INCEPTION V1 (GoogLeNet) VGG NET (Runner up)	6.67 7.3
2015	ResNet	3.57
2016	ResNeXt	4.1
2017	SENet	2.251

AlexNet Architecture

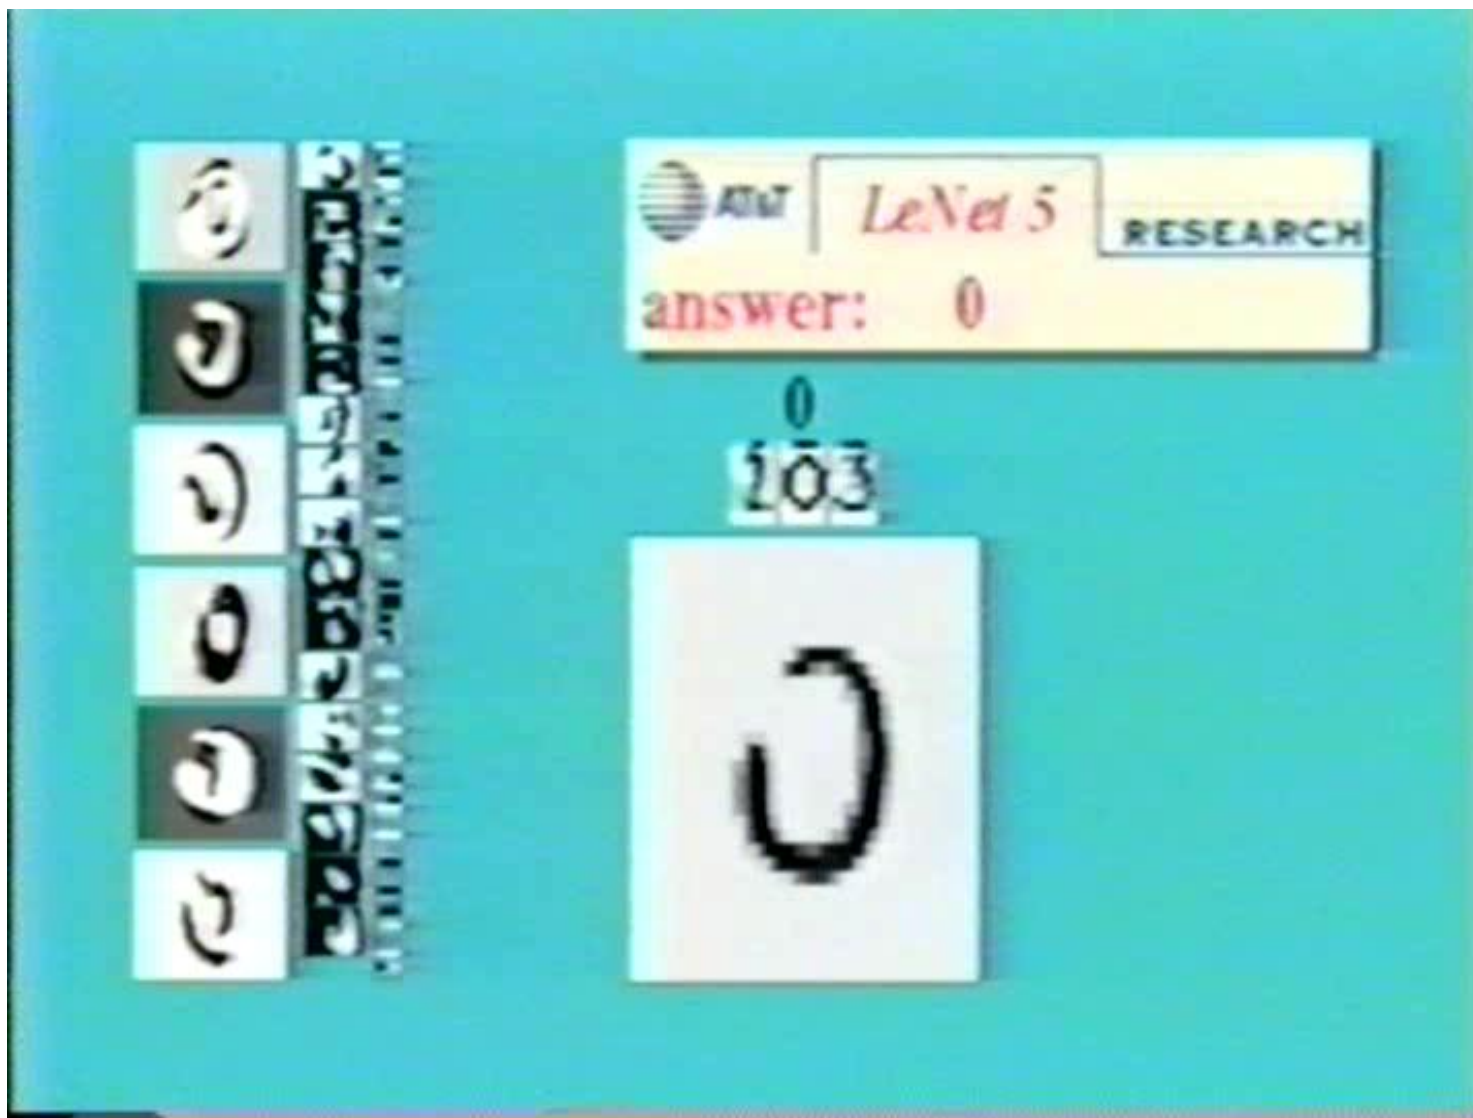


LeNet5 Architecture

- LeNet-5 is the earliest CNN architecture, developed by Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner in 1998.
- It was primarily designed for handwritten digit recognition tasks, particularly recognizing digits in postal codes on letters

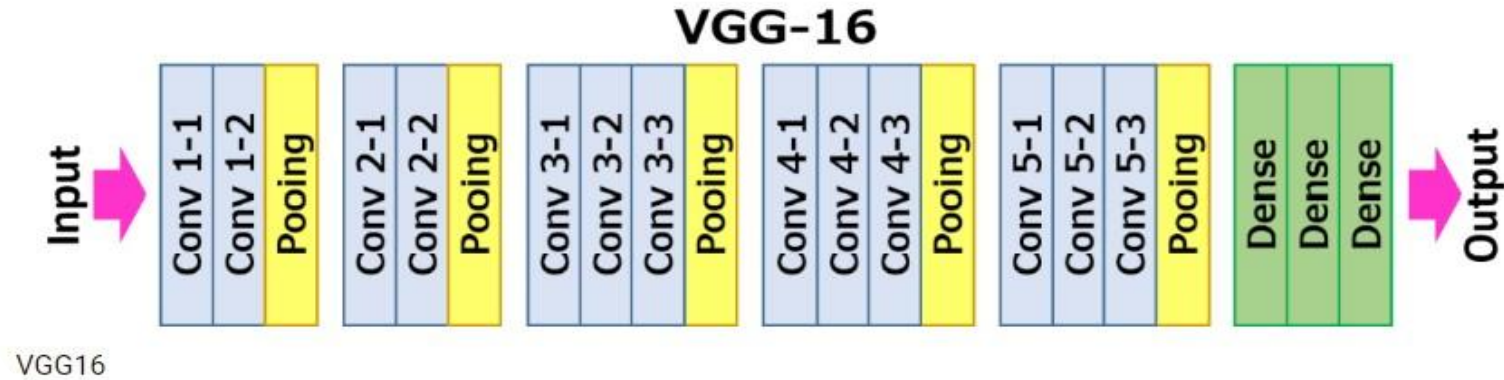


Handwritten digit classification



VGG16/19 Architecture


- VGG-16 model, created by the Visual Geometry Group at the University of Oxford.
- The VGG-16 model is a 16-layer (convolution and fully connected) network built on the ImageNet database, which is built for the purpose of image recognition and classification.

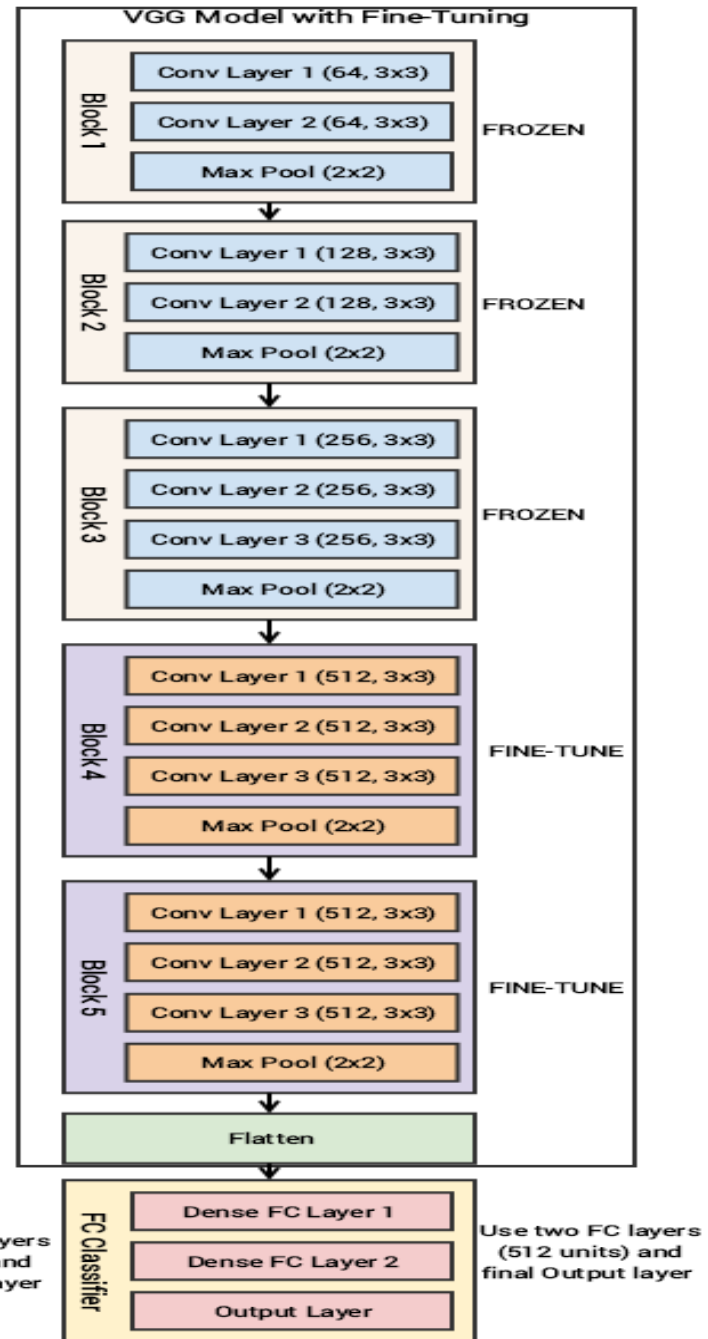
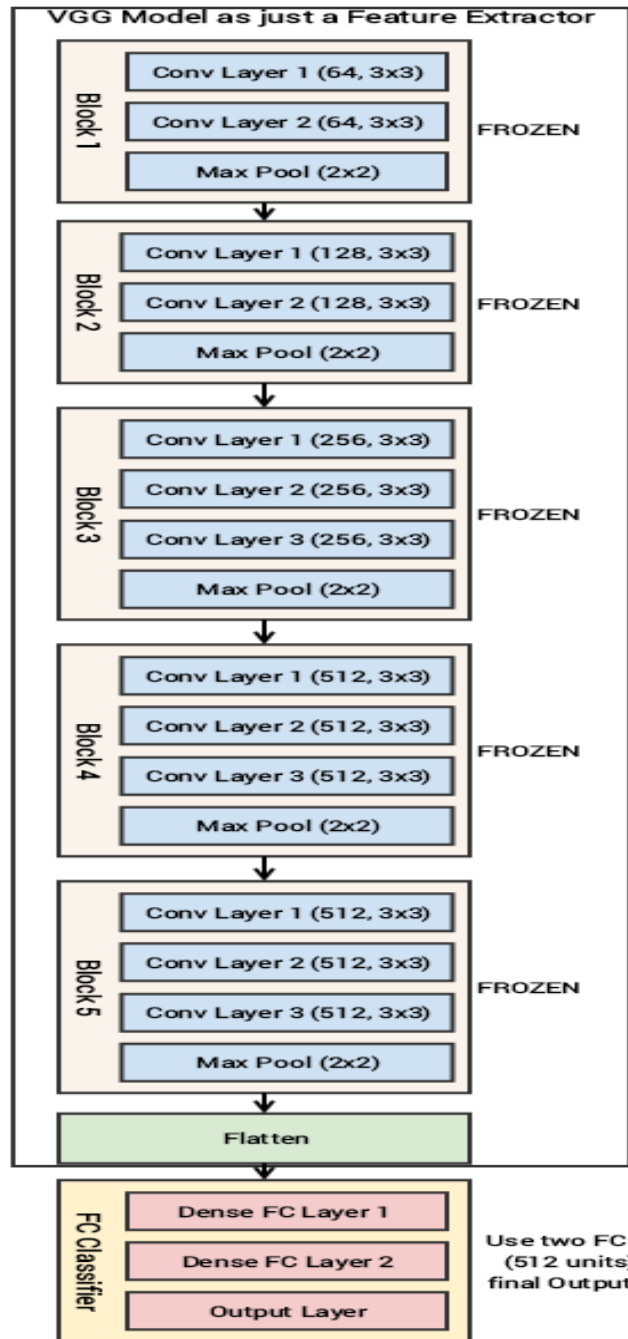



VGG - 16

CONV = 3×3 filter, $s = 1$, same

MAX-POOL = 2×2 , $s = 2$






Problems with Very Deep Networks

- **Initial Belief:**
 - Increasing the number of layers in a neural network would consistently improve accuracy.
- **Practical Issues:**
 - No generalization
 - Vanishing gradient
 - Exploding gradients
 - Hampered training and limited performance

Vanishing
Gradient

Exploding
Gradients



Vanishing/Exploding Gradients Problems:

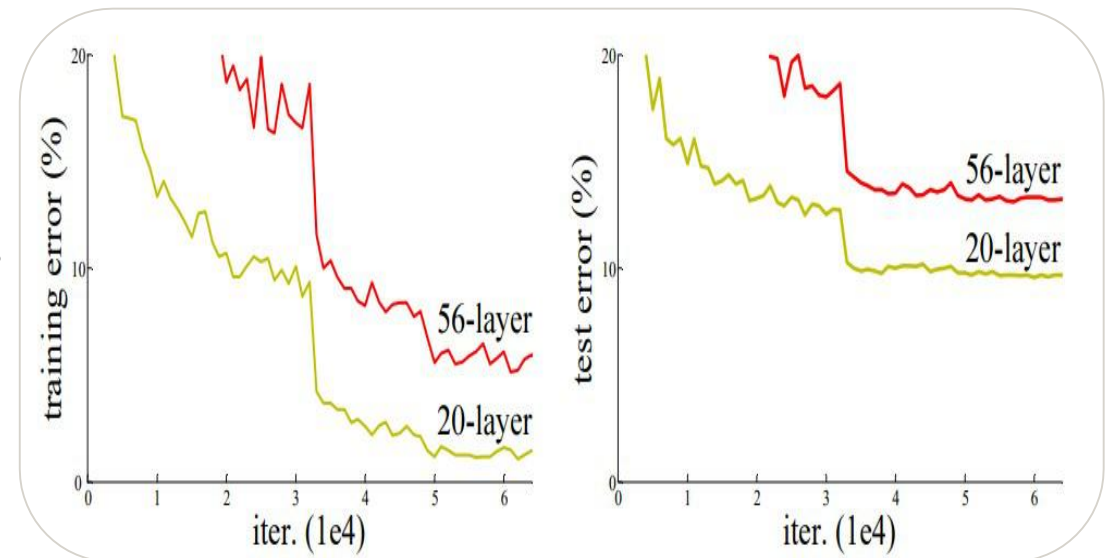
Vanishing Gradient Problem:

- During backpropagation, gradients can become extremely small.
- Lower layers' weights remain unchanged, hindering convergence.

Exploding Gradients Problem:

- During backpropagation, gradients grow larger, leading to unstable training.
- Results in excessively large weight updates.
- Often observed in recurrent neural networks and deep architectures.

It makes learning difficult, particularly in deeper layers.



Comparison of 20-layer vs 56-layer architecture

Solution

Residual Learning

Skip Connections:

- A shortcut connection in a neural network that bypasses some layers and directly feeds the output of an earlier layer to a later layer.
- Helps bypass layers that may degrade performance, effectively regularizing the network.

Residual Blocks:


- Introduce skip connections to address vanishing/exploding gradient issues.
- Form the basic building blocks of ResNet.

Residual Mapping:

- Instead of learning the underlying mapping $H(x)$, the network learns the residual $F(x)$ where $F(x) := H(x) - x$
- Allows the network to fit the function $H(x) := F(x) + x$.

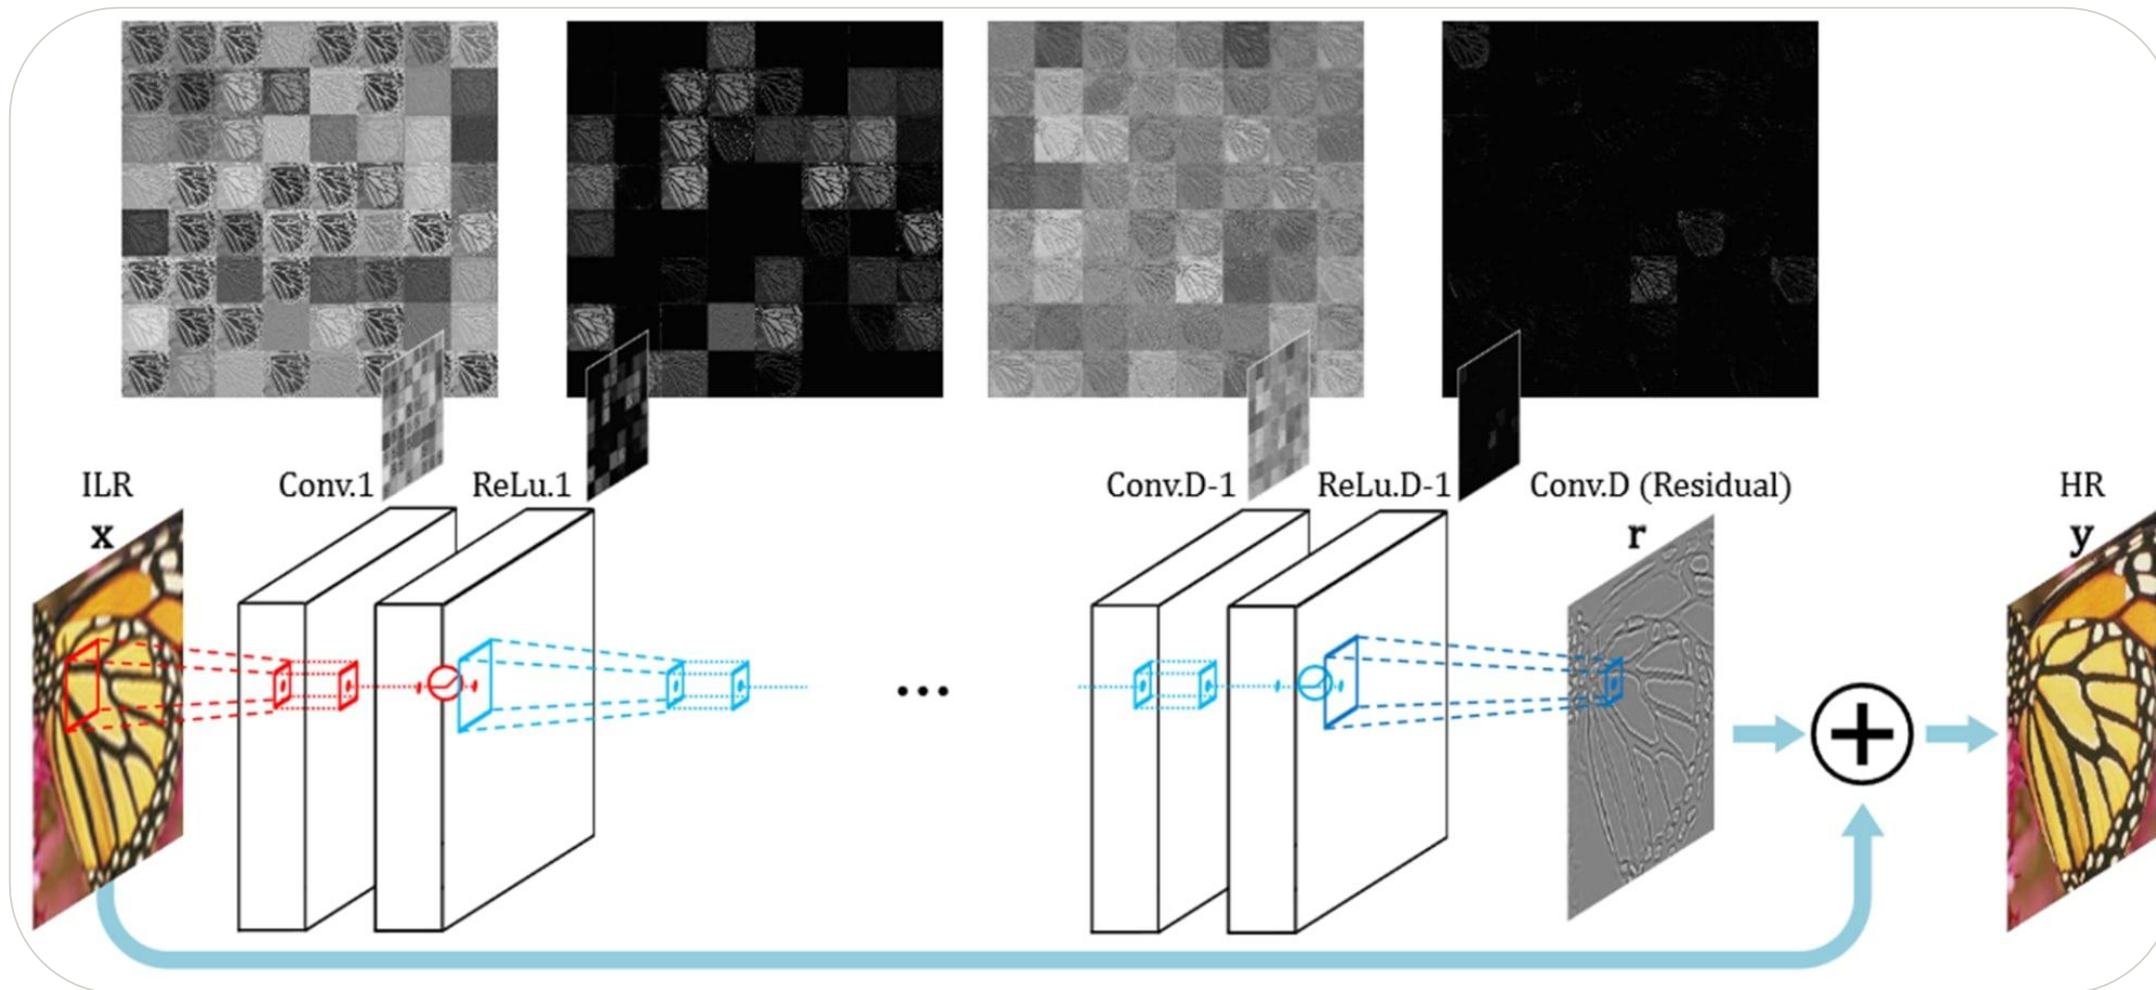
Advantages:

- Mitigates vanishing/exploding gradient problems.
- Enables training of very deep neural networks.
- Improves performance and generalization.



A Residual Block in a deep Residual Network. Here the Residual Connection skips two layers

Residual Learning




ResNet

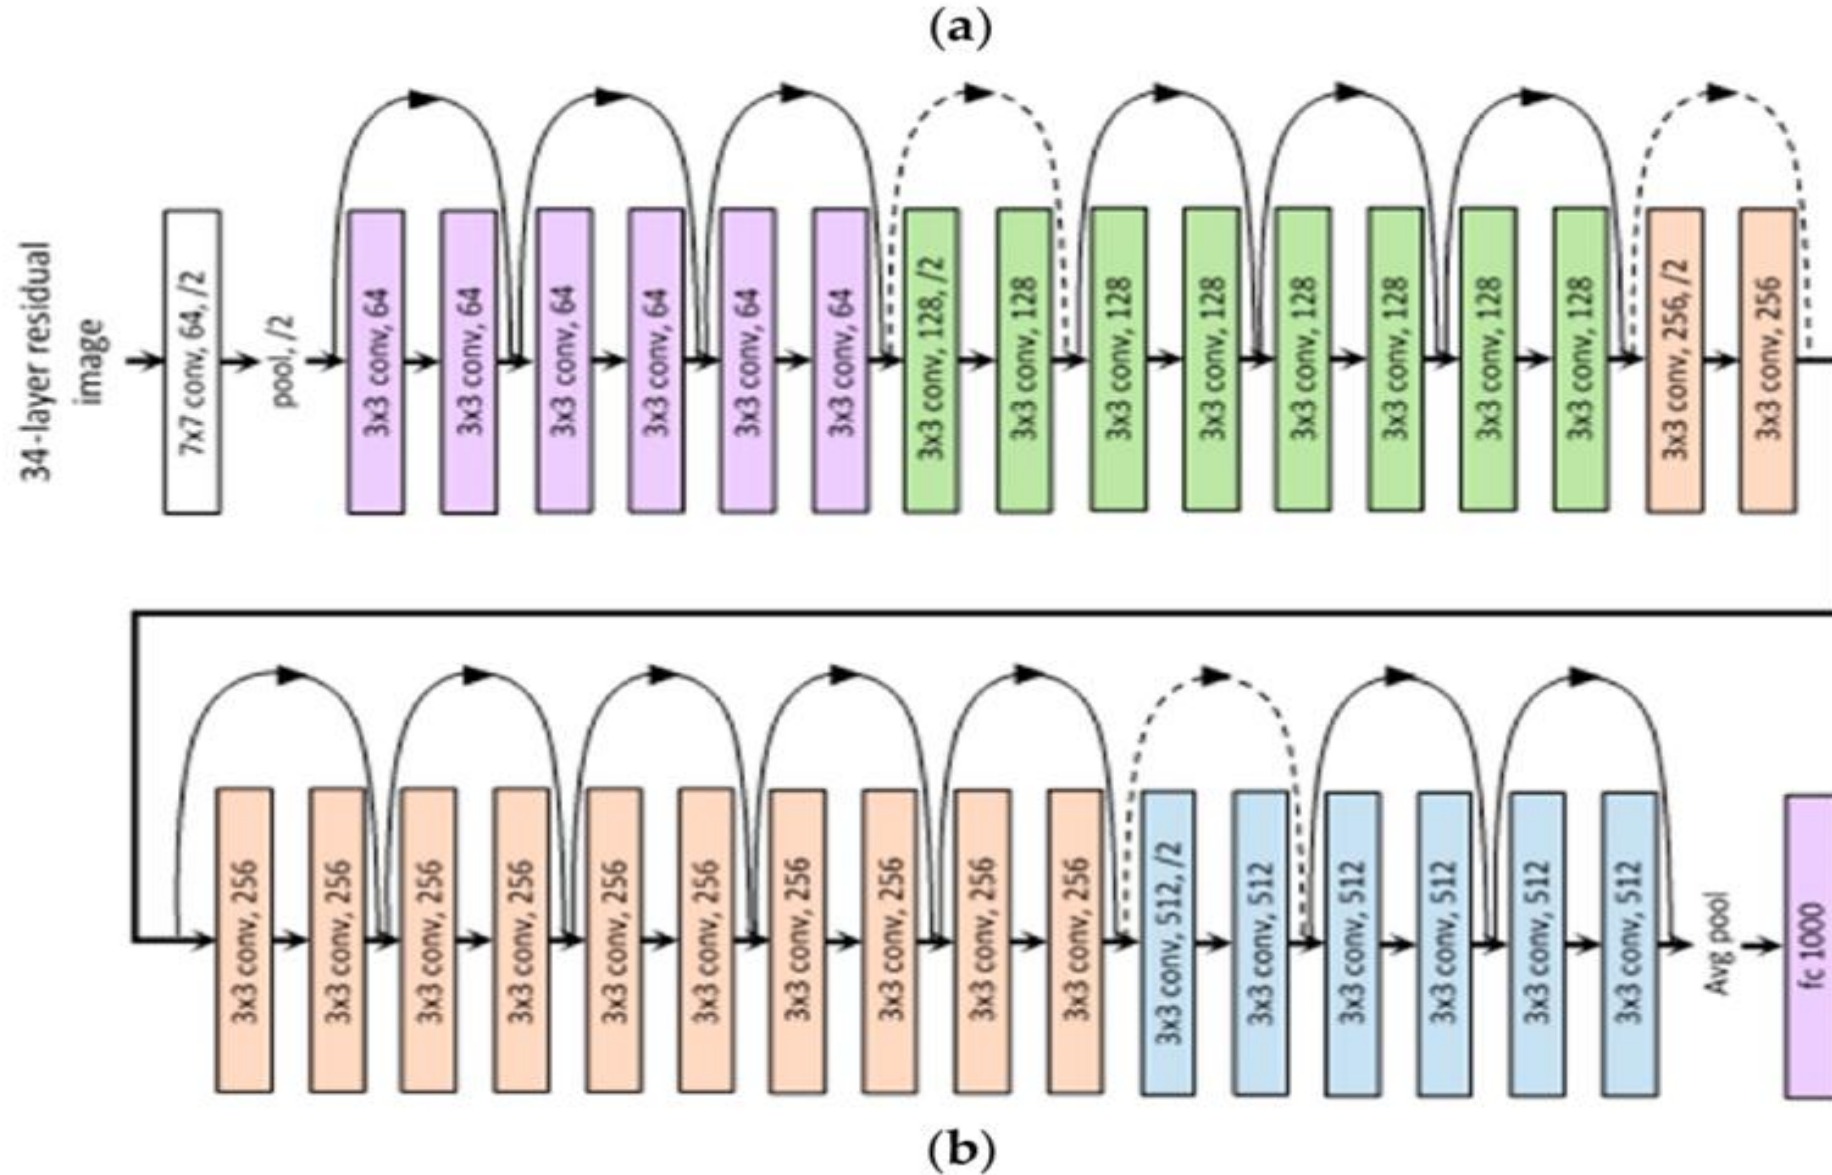
- ResNet, introduced in 2015 by Microsoft Research, proposed a new architecture called Residual Network.
- A ResNet is a type of Convolutional Neural Network (CNN) that stacks residual blocks to form deep architectures.
- ResNet first introduced the concept of skip connection.
- Winner of the ImageNet Challenge in 2015 with an error rate of 3.57%.
- ResNet50 is a variant of ResNet model which has 48 Convolution layers along with 1 MaxPool and 1 Average Pool layer

Variants of ResNet architecture

- Resnet-18, Resnet-34, Resnet-50, Resnet-101, Resnet- 152.
The number after all the model is the number of layers in the model.



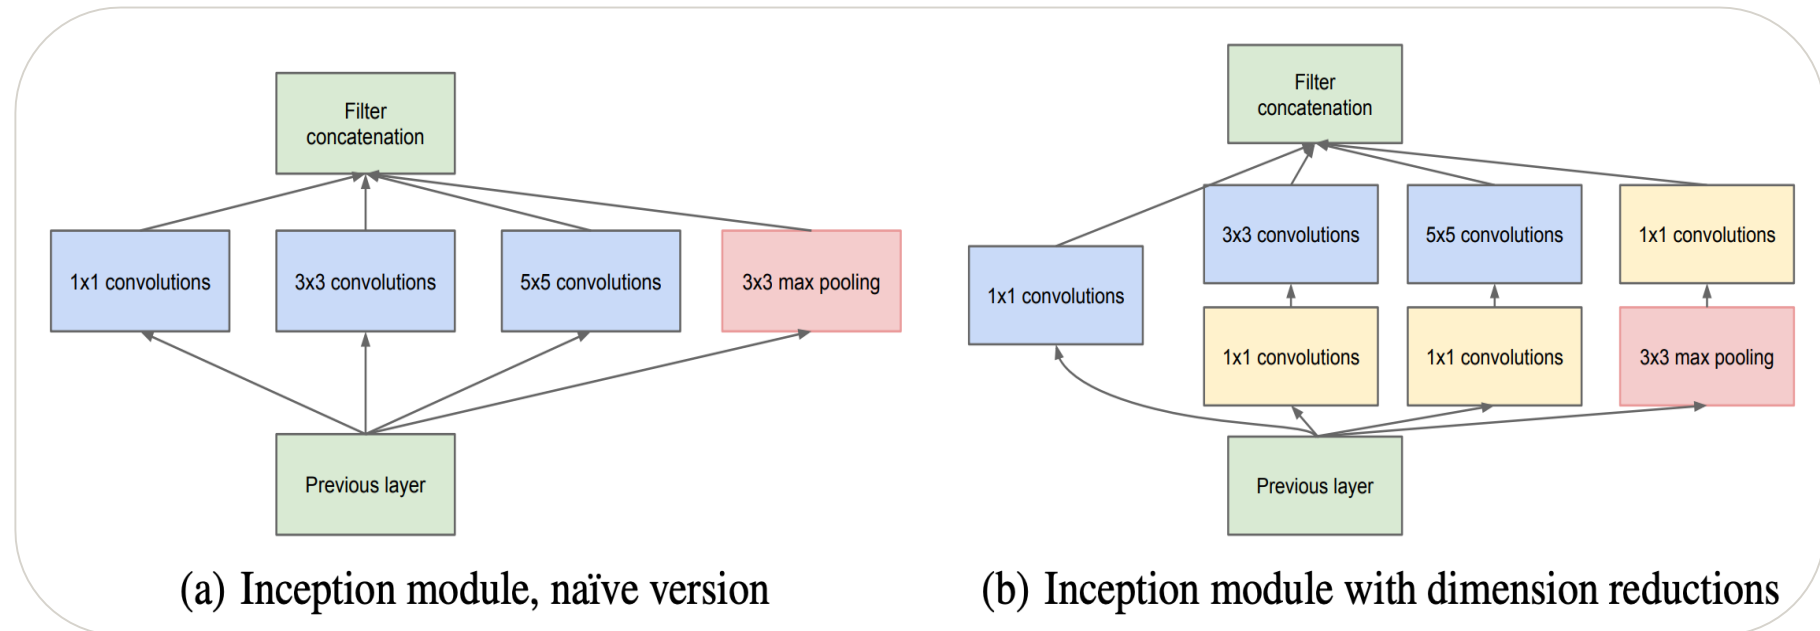

ResNet-34 Layered architecture



Inception Architecture

Inception Module:

- Utilizes multiple convolutional filters (1x1, 3x3, 5x5) and pooling operations within a same module.
- While some networks like VGG16 focus only on 3x3 or LeNet5 on 5x5, Inception makes sure to grab all kinds of features.
- By using various filter sizes, Inception can pick up both small and big details in the data.
- Effectively captures information at different spatial resolutions.

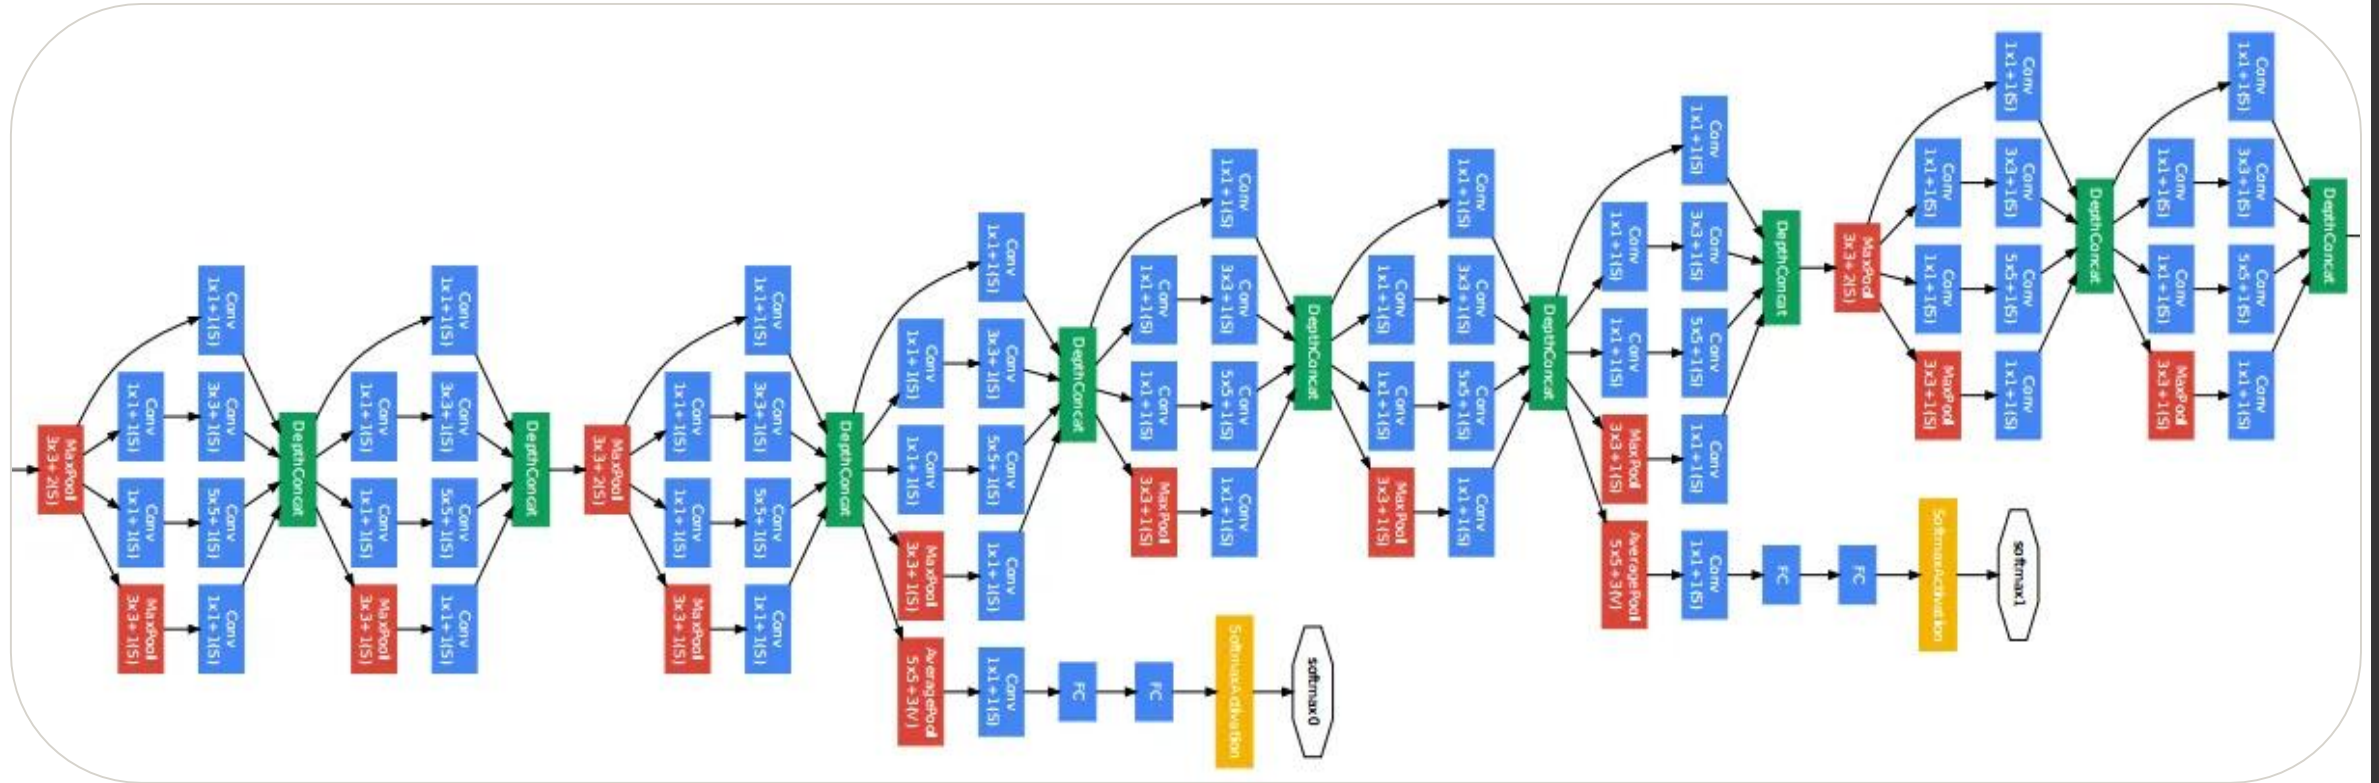


Inception Pre-trained Models:

- Inception-v1 (GoogLeNet), Inception-v2, Inception-v3, Inception-v4.


Notable Achievements:

- Inception-v1 (GoogLeNet) won the 2014 ImageNet Challenge with a top-5 error rate of 6.67%.



GoogLeNet, 2014


The Problem of Computational Cost



$$28 \times 28 \times \underline{192}$$

→
CONV
 5×5 ,
same,
32

32 filters.



$$\underline{28 \times 28 \times 32}$$

filters one $5 \times 5 \times 192$


$$\underline{28 \times 28 \times 32} \times \underline{5 \times 5 \times 192} = \underline{120M.}$$

Solution

Less Parameters means Less Computational Cost.

- Add 1×1 Conv before 3×3
- Add 1×1 Conv before 5×5
- And Add 1×1 Conv after the 3×3 MaxPool layer.

Using 1×1 convolution



1x1 Convolution



- A 1x1 convolution applies a single 1x1 filter to each pixel in the input volume.
- It processes each pixel individually but across all channels (depth), combining the information from different channels.

Purpose:

- Reduces the number of channels while retaining spatial dimensions.
- Enables efficient dimensionality reduction and computational cost savings.

Applications:

- **ResNet:** Used in bottleneck blocks for efficiency.
- **MobileNet:** Part of depthwise separable convolutions.
- **Inception Modules:** Reduces dimensions before expensive convolutions.



EfficientNet

Introduced in 2019 by a team of researchers at Google AI.


- The most powerful CNN architecture.
- EfficientNet is built upon a concept called compound scaling.
- **Compound scaling** optimizes model depth, width, and resolution for optimal efficiency.

Applications:

- Image classification, object detection, semantic segmentation

EfficientNet Variants:

- **EfficientNet B0-B7:** A family of EfficientNet models with varying complexities.
- **B0: Most lightweight** (5.3 million parameters).
- **B7: Most complex** (6.1 billion parameters).



MobileNet

Developed by Google researchers.

Purpose:

- Designed for mobile and embedded vision applications.
- Focuses on efficient, lightweight models suitable for devices with limited computational resources.

Key Features:

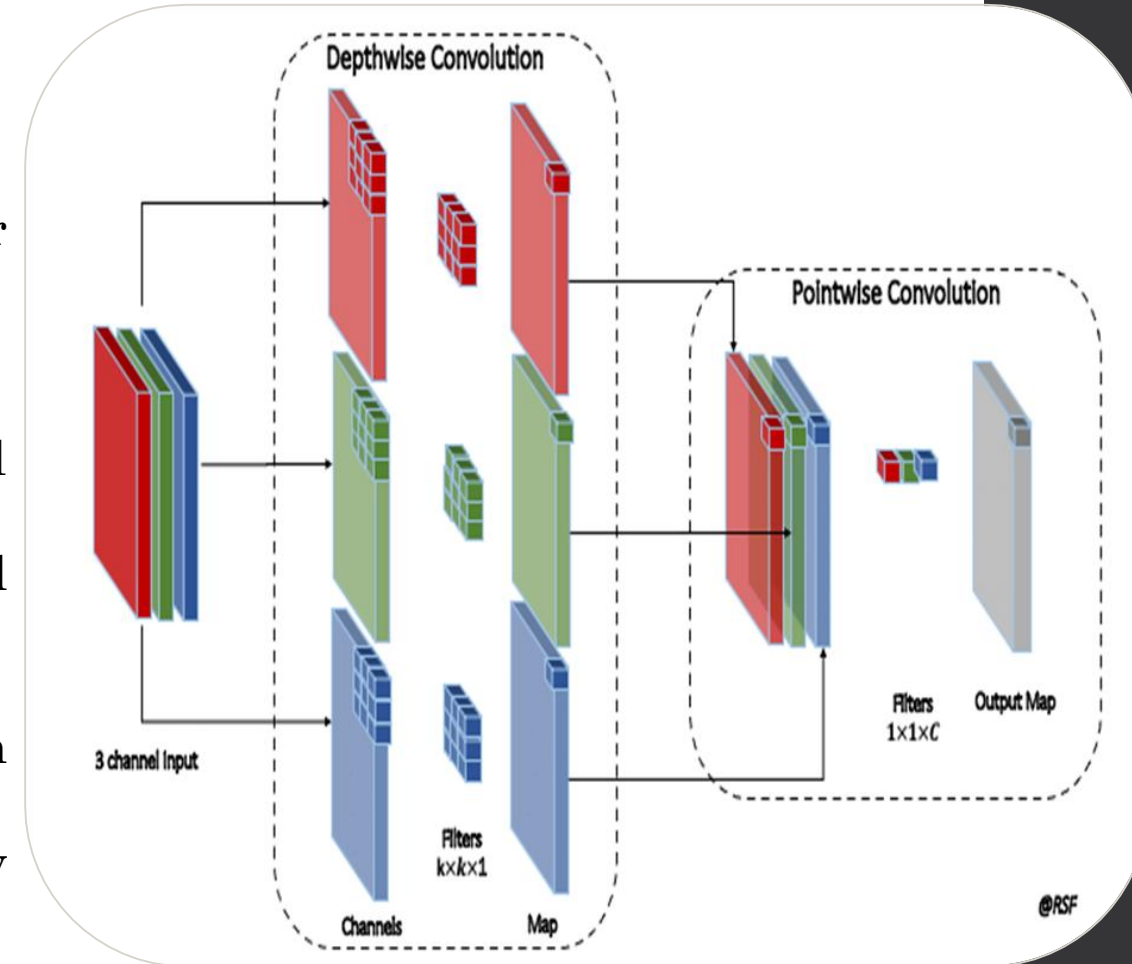
- Reduces computational cost and model size.
- Fewer parameters compared to traditional convolutional networks.
- Maintains competitive accuracy with optimized speed and efficiency.

Applications:

- Real-time object detection and image classification on mobile devices.
- Deployment in IoT devices and augmented reality applications.

Pre-trained Models:

- MobileNetV1, MobileNetV2, MobileNetV3.
- Pre-trained on ImageNet, available for transfer learning.



Applications of transfer Learning



Applications of transfer Learning

Image Classification

A core application of transfer learning in computer vision.

Pre-trained Models

- Leverage powerful models like ResNet, VGG, and Inception.
- Trained on massive datasets like ImageNet.
- Fine-tune models for specific domains.

Applications:

- For examples, Identifying species in wildlife photography or diagnosing medical conditions from imaging data.

Object Detection:



Detect and localize objects in images or videos.

Pre-trained Models:

- Utilize pre-trained models like YOLO, Faster R-CNN, SSD for feature extraction.
- Add layers for bounding box prediction and class identification.

Applications:

- Pedestrian detection for self-driving cars.



Applications of transfer Learning

Image Segmentation

Segmenting images into distinct regions corresponding to objects or parts of objects.

Pre-trained Models:

- U-Net, DeepLab, FCN.

Applications with Transfer Learning

- **Medical Imaging:** Identify tumors or other abnormalities.
- **Autonomous Vehicles:** Differentiate between roads, sidewalks, and vehicles.



Face Recognition:


Identify and verify faces in images or videos.

Pre-trained Models:

- Utilize pre-trained models like FaceNet, VGGFace for feature extraction.
- Add layers for face identification and verification.

Applications:

- Security systems for access control.
- Social media tagging.
- User authentication for devices and apps.



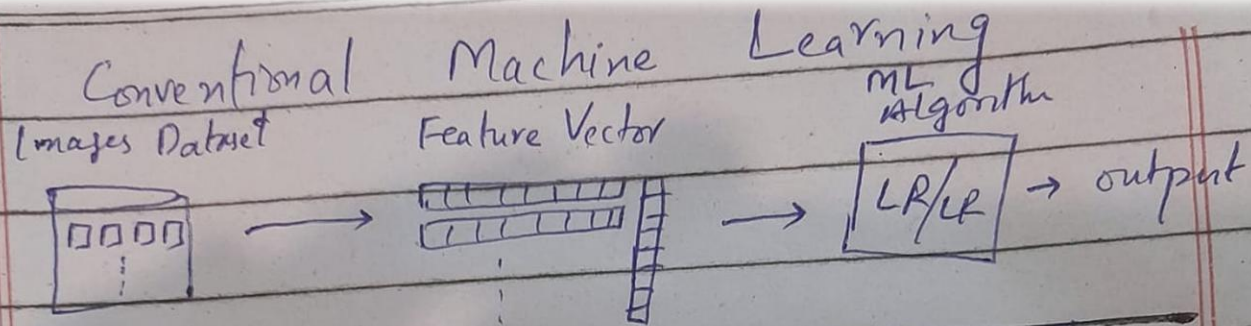


References:


- https://www.researchgate.net/figure/Change-in-the-number-of-model-parameters-and-training-time-with-increased-number-of_fig3_335865710
- <https://blog.devgenius.io/all-you-need-to-know-about-transfer-learning-4a7e3cbaf1dd>
- <https://medium.com/@nutanbhogendrasharma/transfer-learning-using-feature-extraction-in-deep-learning-afd97380c96c>
- https://dev.mrdbourke.com/tensorflow-deep-learning/04_transfer_learning_in_tensorflow_part_1_feature_extraction/
- <https://encord.com/glossary/pre-trained-model-definition/>
- <https://en.wikipedia.org/wiki/ImageNet>
- <https://blog.mturk.com/tutorial-how-to-label-thousands-of-images-using-the-crowd-bea164ccbefc>
- <https://medium.com/@siddheshb008/alexnet-architecture-explained-b6240c528bd5>
- 62.3 million learnable parameters.
- <https://www.philschmid.de/getting-started-with-cnn-by-calculating-lenet-layer-manually>
- <https://medium.com/@mygreatlearning/everything-you-need-to-know-about-vgg16-7315defb5918>
- <https://www.geeksforgeeks.org/residual-networks-resnet-deep-learning/>
- <https://viso.ai/deep-learning/resnet-residual-neural-network/>
- https://www.researchgate.net/figure/A-schematic-view-of-ResNet-architecture-15-decomposed-into-three-blocks-embedding_fig1_333475917

References:

- <https://medium.com/@siddheshb008/resnet-architecture-explained-47309ea9283d>
- <https://velog.io/@snoop2head/Going-Deeper-with-Convolution-GoogleNet-Inception>
- <https://towardsdatascience.com/an-intuitive-guide-to-deep-network-architectures-65fdc477db41>
- <https://www.youtube.com/watch?app=desktop&v=pf-HUcqdCr4>
- https://www.baeldung.com/wp-content/uploads/sites/4/2020/06/3D_1D_cropped.gif
- <https://arxiv.org/abs/1905.11946>
- <https://deeplobe.ai/image-segmentation-the-most-interesting-applications/>
- <https://medium.com/analytics-vidhya/how-facial-recognition-systems-work-edcbb227e614>



Transfer Learning



Thank You